

C-OWPF Codebase Reference

A Complete, Detailed Summary of the Python Implementation
Optimal Water Flow (OWF) and Coupled Optimal Water–Power Flow (C-OWPF)

Companion documentation to IEEE Access-2024-18604

August 7, 2026

Abstract

This document is a complete, self-contained reference for the Python implementation of the C-OWPF solver. Part I gives the architecture and the core data structures. Part II explains, in detail, *how* each optimization problem is solved — the decoupled Optimal Water Flow (OWF, problem P3) and the coupled Optimal Water–Power Flow (C-OWPF, problem P2) — naming the exact function call chain, the variables passed in, and the outputs produced at each stage. Part III is an auto-generated per-module API reference documenting every function and class (168 functions, 20 classes across 30 modules): its signature, its inputs, and its outputs. The reference is generated directly from the source by an AST walk, so it is guaranteed to match the code.

Contents

I	Architecture and Core Data Structures	3
1	Package layout	3
2	The solver stack	3
3	Core data structures	3
II	How the Problems Are Solved	5
4	Decoupled Optimal Water Flow (OWF, problem P3)	5
4.1	Entry point and inputs	5
4.2	Building the problem: <code>owf.network.setup</code>	5
4.3	The successive-linearization loop: <code>owf.solver.solve_owf</code>	5
4.4	The four modes (how the schedule is obtained)	6
4.5	Validation: <code>owf.validation.validate_schedule</code>	6
5	Coupled Optimal Water–Power Flow (C-OWPF, problem P2)	6
5.1	Entry point and inputs	6
5.2	One coupled solve: <code>coupled.coupled_lp.solve_coupled</code>	6
5.3	The voltage-aware schedule search: <code>coupled.schedule.optimize_coupled_schedule</code>	7
5.4	Validation: <code>coupled.validation.validate_coupled</code>	7
5.5	Related coupled drivers	7
III	Complete Per-Module API Reference	8
6	Entry points	8
6.1	<code>main.owf.py</code>	8
6.2	<code>app.py</code>	9
6.3	<code>guide.py</code>	9

7	owf – Water Distribution Network & OWF solver	9
7.1	owf/___init___py	9
7.2	owf/config.py	9
7.3	owf/epanet_io.py	10
7.4	owf/connection_matrices.py	11
7.5	owf/network.py	12
7.6	owf/initial_values.py	13
7.7	owf/linearization.py	13
7.8	owf/constraints.py	14
7.9	owf/solver.py	16
7.10	owf/warmstart.py	17
7.11	owf/validation.py	19
7.12	owf/plots.py	19
7.13	owf/netmap.py	20
8	pdn – Power Distribution Network models	20
8.1	pdn/___init___py	20
8.2	pdn/feeders.py	20
8.3	pdn/feeders_paper.py	21
8.4	pdn/network.py	21
8.5	pdn/lindistflow.py	21
8.6	pdn/powerflow.py	22
8.7	pdn/opf.py	23
9	coupled – Coupled C-OWPF solver	23
9.1	coupled/___init___py	23
9.2	coupled/config.py	23
9.3	coupled/coupled_lp.py	24
9.4	coupled/runner.py	25
9.5	coupled/schedule.py	25
9.6	coupled/validation.py	26
10	ui – Streamlit application	26
10.1	ui/___init___py	26
10.2	ui/theme.py	26
10.3	ui/landing.py	27
10.4	ui/water_section.py	27
10.5	ui/power_section.py	28
10.6	ui/coupled_section.py	28
10.7	ui/pdn_plots.py	28
11	tests	29
11.1	tests/test_coupled.py	29
11.2	tests/test_eightnode.py	30
11.3	tests/test_more_networks.py	30
11.4	tests/test_optimize_schedule.py	31
11.5	tests/test_plots.py	32

Part I. Architecture and Core Data Structures

1 Package layout

The implementation lives under `Python/` and is organized into four packages plus a command-line entry point and a Streamlit application:

Package	Responsibility
<code>owf/</code>	The water side. Parses EPANET <code>.inp</code> files, builds the WDN problem data and incidence matrices, forms the successive-linearization coefficients, assembles the CVXPY constraints/objective, runs the MILP loop, warm-starts and searches pump schedules, and validates the result against EPANET.
<code>pdn/</code>	The power side. Feeder datasets, the linear LinDistFlow/Kekatos voltage model, the nonlinear Z-bus power-flow validator, and a standalone reactive-power OPF.
<code>coupled/</code>	The joint C-OWPF. Augments the water model with the feeder voltage model, PV reactive support, and a linearized loss term; a voltage-aware pump-schedule search; and a two-simulator (EPANET + Z-bus) validator.
<code>ui/</code>	The Streamlit application (Home / Water / Power / Coupled / Guide tabs).
<code>main_owf.py</code>	Command-line driver for the decoupled OWF (<code>run_case</code>).
<code>app.py</code>	Streamlit entry point.

2 The solver stack

Both problems are nonconvex (Hazen–Williams head loss $\sim q^{1.852}$, FSP pump power, and the AC power flow). They are solved by **successive linear approximation**: the nonlinear terms are linearized about a current operating point, producing a mixed-integer linear/quadratic program that is solved by **HiGHS** (default) or **SCIP** (fallback) through **CVXPY**; the model is then relinearized about the new solution and re-solved until the iterate stops moving. Only the pump on/off variables are integer; the feeder (LinDistFlow + PV reactive) adds no new integers.

3 Core data structures

The problem data and results flow through a small number of dataclasses. Their full field lists appear in Part III; the summary below orients the reader.

Type	Module	Role
RawNetwork	owf/epanet_io	Static data read from an EPANET <code>.inp</code> : nodes, links, from/to, resistances, pump indices, tanks, reservoirs, junctions (0-based).
Matrices	owf/connection_matrices	Indices / selection matrices ($\Pi, \Pi', \Lambda, \Theta, T, K, \Omega, \Delta, \dots$).
WDN	owf/network	The fully assembled water problem: matrices, pump/tank params, bounds, price, availability, EPANET ground truth. Passed to every water solve.
LinPoint	owf/linearization	Frozen linearization coefficients for one MILP solve (pipe/pump/power tangents).
OWFResult	owf/solver	Output of a water solve: flows, heads, on/off, pump power, cost, per-iteration error/objective, convergence, solver.
LinDistModel	pdn/lindistflow	Affine feeder voltage model $v^2 = R(-p) + X(-q) + V_k$ (all pu).
PDN	pdn/network	A feeder ready for coupling: its LinDistModel , PV/solar data, and voltage limits.
CoupledConfig	coupled/config	How a WDN and feeder are co-optimized (feeder key, pump→bus map, PV sizing, $v_{\min}/v_{\max}$, PF, loss on/off).
CoupledResult	coupled/coupled_opt	Output of a coupled solve: everything in OWFResult plus bus voltages, PV active/reactive, loss cost, total cost, voltage violation.

Part II. How the Problems Are Solved

4 Decoupled Optimal Water Flow (OWF, problem P3)

4.1 Entry point and inputs

A water case is run through `main_owf.run_case`:

```
run_case(net, mode, price, horizon, plot, outdir, verbose, solver='HIGHS')
-> (CaseResult, wdn, result)
```

Inputs: `net` (network id, e.g. 8 = 8-node, 97 = Net3, 126 = BWSN); `mode` $\in$ {direct, warmstart, optimize, epanet}; `price` (1 = time-of-use, 0 = flat); `horizon` (hours, or None for the network default); `solver`. **Output:** a `CaseResult` (costs, savings, EPANET-replay errors, timing), plus the WDN object and the `OWFResult`.

4.2 Building the problem: `owf.network.setup`

`run_case` first constructs a `SolverConfig` and calls `setup(config) -> WDN`, which: (i) reads the `.inp` via `epanet.io.read_inp -> RawNetwork`; (ii) runs an EPANET extended-period simulation (`run.epanet`) to obtain ground-truth flows/heads; (iii) builds the incidence matrices (`connection_matrices.build_matrices -> Matrices`); (iv) derives pump, tank, bound, price, and availability data (`_pump_params`, `_tank_params`, `_bounds`, `_price`, `_resolve_availability`). The returned WDN carries all of this and is the single object passed to every solver.

4.3 The successive-linearization loop: `owf.solver.solve_owf`

```
solve_owf(wdn, lin_override=None, eps_override=None) -> OWFResult
```

Steps:

1. **Initial point.** `initial_values.initial_point` returns a `LinPoint` (linearization tangents), the stacked iterate `eps` = [Heads; Flows; OnOff], and an initial `OnOff`. It can be seeded from EPANET flows/heads.
2. **Model.** `_build_model(wdn) -> Model` declares the CVXPY variables — nodal `Heads`, link `Flows`, binary `OnOff`, `Ppump`, and (in soft mode) bound slacks — and the linearization *parameters*.
3. **Constraints and objective.** `constraints.build_constraints(model, wdn)` assembles mass balance, reservoir/tank/junction head relations, the linearized pipe head loss (Π -based), the big- M pump curve, pump flow/power, and availability; `constraints.objective` sets $\min \sum_t \pi_t \sum_p P_{p,t}^{\text{pump}} / 1000$.
4. **Iterate.** Each iteration `_set_params(model, lin)` injects the current `LinPoint`; the MILP is solved by HiGHS/SCIP; `linearization.linearize(flows, ...)` recomputes the tangents about the new flows; `error_norm` measures the change in `eps`. The loop stops at `tolerance` or `max_iter`.

Output — `OWFResult`: `flows` ($L \times T$), `heads` ($N \times T$), `onoff` ($P \times T$), `ppump` (linear) and `ppump_true` (nonlinear), scalar `cost`, per-iteration `errors/objective`, `converged`, `n_iter`, `solver`.


4.4 The four modes (how the schedule is obtained)

Mode	Driver (owf/)	What it does
epanet	warmstart.solve_from_epanet	Reproduces EPANET's rule-based operation (cost reference).
direct	solver.solve_owf	One successive-linearization run from the default initial point (radial networks).
warmstart	warmstart.solve_warmstart / solve_multistart	Converges from several EPANET-seeded schedules; robust on looped networks.
optimize	warmstart.optimize_schedule	Searches pump on/off (EPANET baseline → price-quantile candidates → fix-and-converge → 1-opt polish), returns savings vs EPANET.

Supporting routines: `epanet_default_onoff` (pump on/off implied by EPANET), `price_threshold_schedules` / `duty_preserving_schedule` (candidate schedules), `solve_fixed_schedule` / `_converge_fixed` (pin a schedule and converge the continuous problem), `true_energy_cost` (cost from the *nonlinear* FSP power).

4.5 Validation: `owf.validation.validate_schedule`

The optimized `onoff` is imposed back in EPANET (`epanet.io.simulate_with_schedule`) and the true hydraulics are compared to the model, yielding the maximum head/flow deviation and the minimum nodal pressure (feasibility). These populate `CaseResult.max_dhead` and `min_pressure`.



5 Coupled Optimal Water–Power Flow (C-OWPF, problem P2)

5.1 Entry point and inputs

The coupled problem is set up and solved through the `coupled/` package:

```

setup(net_num, cc, time=None, price_choice=1, solver='HIGHS') -> (WDN, PDN)
optimize_coupled_schedule(wdn, pdn, cc, ...) -> (CoupledResult, info)

```

Inputs: `net_num` (water network), a `CoupledConfig` `cc` (feeder key, pump→bus mapping, PV sizing $S_{pv} = k P_{pv}$, $v_{\min}/v_{\max}$, pump power factor, loss term on/off). `setup` builds the WDN (as for OWF) and the PDN via `pdn.network.PDN.build(feeder_key, ...)`, which wraps the affine `lindistflow.build_lindist` → `LinDistModel` together with the feeder's PV/solar and voltage limits.

5.2 One coupled solve: `coupled.coupled_lp.solve_coupled`

```

build_coupled_problem(wdn, pdn, cc, trust_region=None) -> (problem, model, ctx)
solve_coupled(wdn, pdn, cc, lin_override=None, eps_override=None,
               trust_region=None) -> CoupledResult

```

`build_coupled_problem` takes the water `Model` and augments it with the power side:

- **Coupling.** The pump electrical power enters the feeder as a bus active load at the mapped bus; its reactive part is $q^{\text{pump}} = P_{\text{pu}}^{\text{pump}} \tan(\arccos \text{PF})$.
- **Voltages.** $v^2 = R(-p_{\text{net}}) + X(-q_{\text{net}}) + V_k$ (`LinDistFlow`), with voltage-dependent shunt capacitors folded into R, X, V_k .
- **PV reactive.** `_pv_capability` fixes PV active output and gives the exact reactive capability (octagon); `PV VAR` is dispatched to hold voltage.

- **Loss.** The branch-flow loss $\sum_{\ell} r_{\ell}(P_{\ell}^2 + Q_{\ell}^2)$ is linearized each iteration via the subtree matrix `branch_flow_matrix`; the priced loss is added to the objective (paper 33d).
- **Voltage limits.** Enforced hard, or via the water side’s soft-slack feasibility device.
- **Trust region (optional).** `trust_region=(z0, K)` caps the pump on/off Hamming distance from an incumbent `z0` to K flips.

The successive-linearization loop then proceeds exactly as in the water solve (relinearize $\rightarrow$ re-solve), and `_snapshot` reads the converged variables. **Output** — `CoupledResult`: `flows`, `heads`, `onoff`; `voltage` and `v2` ($N \times T$, pu); `pv_p`, `pv_q`; `energy_cost`, `loss_cost`, `total_cost`, `loss_kw`; `v_violation`; `water_max_slack`.

5.3 The voltage-aware schedule search: `coupled.schedule.optimize_coupled_schedule`


Unlike the water-only search, every candidate schedule is scored on the *coupled* solve, so a schedule that is cheap for water but causes an undervoltage is rejected. Stages:

1. **Baseline.** EPANET’s own schedule, evaluated by `runner.solve_coupled_schedule` (fix `onoff`, warm-start from EPANET, converge the coupled LP).
2. **Candidates.** Price-quantile / load-shift schedules (`price_threshold_schedules`).
3. **Warm-started trust-region MILP.** Free binaries, warm-started at the incumbent’s linearization, Hamming cap $K = \max(2, \lfloor n_p T / 4 \rfloor)$ — the MATLAB-style free-binary MILP made reliable.
4. **1-opt polish.** Single pump-hour flips, greediest-price first.

Every candidate is ranked by `_score`: *coupled-feasible first* (water head slack and voltage violation within tolerance), then lowest true total cost (nonlinear pump energy + priced loss). A final clean re-solve of the winner *and* the EPANET baseline guarantees the coupled result is never worse than the decoupled one.

5.4 Validation: `coupled.validation.validate_coupled`

The winning solution is cross-checked against *both* exact simulators: the water schedule is replayed in EPANET (head/flow error, pressures), and the bus injections are replayed in the nonlinear Z-bus solver (`powerflow.zbus.solve`) for the true voltages and true $\sum |I|^2 r$ loss (`coupled_loss_kwh`).



5.5 Related coupled drivers

`runner.solve_coupled_schedule` (fixed-schedule coupled solve, used as the scoring oracle); `runner.solve_coupled_epanet` (coupled solve on EPANET’s own schedule); `coupled_lp.default_pump_buses` (auto-place pumps at the electrically weakest buses); `opf.solve_pdn_opf` (the standalone power-only reactive OPF used by the Power tab).

Part III. Complete Per-Module API Reference

Auto-generated from the source by an AST walk. For each function: its signature (inputs) and return annotation (outputs) followed by its docstring. For each class: its fields and methods.

6 Entry points

6.1 main_owf.py

Module purpose. Entry point for the FSP Optimal Water Flow (OWF) solver.

`_fallbacks`

```
def _fallbacks(solver: str) -> tuple:
    """Fall back to HiGHS then SCIP (whichever are installed), skipping the primary."""
```

`_warmstart_config`

```
def _warmstart_config(net: int, price: int, horizon, solver: str) -> SolverConfig:
```

`class CaseResult`

```
class CaseResult:
    # fields:
    label: str
    net_num: int
    mode: str
    price: str
    horizon: int
    elapsed: float
    converged: bool
    n_iter: int
    epanet_cost: float = float('nan')
    owf_cost: float = float('nan')
    savings_pct: float = float('nan')
    max_dhead: float = float('nan')
    max_dpumpflow: float = float('nan')
    min_pressure: float = float('nan')
    solver: str = 'HiGHS'
    note: str = ''
    plots: list = field(default_factory=list)
```

`run_case`

```
def run_case(net: int, mode: str, price: int, horizon, plot: bool, outdir: str, verbose: bool, solver: str='HiGHS') -> tuple[Optional[CaseResult], object, object]:
    """Construct and solve one case; returns (CaseResult, wdn, result)."""
```

`comparison_table`

```
def comparison_table(cases: list) -> str:
    """EPANET(rules) vs C-OWPF comparison for one or more cases."""
```

`_ask`

```
def _ask(prompt: str, default: str, valid=None) -> str:
```

`interactive`

```
def interactive() -> None:
```


main

```
def main() -> None:
```

6.2 app.py

Module purpose. C-OWPF Streamlit app – coupled Optimal Water-Power Flow.

6.3 guide.py

Module purpose. Methodology / user guide rendered inside the Streamlit app (Guide tab).

render_guide

```
def render_guide(st) -> None:
```

7 owf – Water Distribution Network & OWF solver

7.1 owf/__init__.py

Module purpose. FSP Optimal Water Flow (OWF) – Python translation of the MATLAB C-OWPF WDN solver.

__getattr__

```
def __getattr__(name):
```

7.2 owf/config.py

Module purpose. Configuration for the FSP Optimal Water Flow (OWF) solver.

class NetworkSpec

```
class NetworkSpec:
    """Everything network-specific needed to build the OWF problem.

    ‘pump_coefficients’ is an optional override: one ‘[h0, r, v]’ row per pump
    giving the pump head-gain curve  $H_{\text{gain}}(q) = h_0 - r * q^{**v}$ . When ‘None’
    (the default) the coefficients are derived from the EPANET pump head curve,
    which keeps the optimizer consistent with EPANET and avoids the value drift
    seen in the hard-coded ‘Prepare_net_WDN.m’ tables."""
    # fields:
    net_num: int
    name: str
    inp_relpath: str
    pump_coefficients: Optional[List[List[float]]] = None
    bypasses: Optional[dict] = None
    pump_availability: Optional[dict] = None
```

method inp_path

```
def inp_path(self) -> Path:
```

class SolverConfig

```
class SolverConfig:
    """Successive-linearization loop and MILP solver settings."""
    # fields:
    net_num: int = 8
    time: Optional[int] = None
    price_choice: int = 1
    choice: int = 1
    tol: float = 0.5
    max_iter: int = 50
    big_m: Optional[float] = None
    mass_balance_warmup: bool = False
```

```

mass_balance_warmup_iters: int = 10
damping: float = 1.0
fixed_schedule: Optional['np.ndarray'] = None
soft_bounds: bool = False
penalty_weight: float = 1000.0
penalty_growth: float = 2.0
penalty_max: float = 10000000.0
feas_tol: float = 0.001
solver: str = 'HIGHS'
fallback_solvers: tuple = ('SCIP',)
verbose: bool = False
data_dir: Path = field(default=DATA_DIR)

```

method spec

```
def spec(self) -> NetworkSpec:
```

available_solvers

```
def available_solvers() -> list[str]:
    """MILP solvers from SOLVER_CHOICES that are actually installed."""
```

7.3 owf/epanet_io.py

Module purpose. EPANET input parsing and hydraulic ground-truth via epyt.

class RawNetwork

```

class RawNetwork:
    """Static network data extracted from an EPANET .inp file (0-based indices)."""
    # fields:
    link_name_id: list[str]
    link_diameter: np.ndarray
    link_length: np.ndarray
    link_roughness: np.ndarray
    link_resistance: np.ndarray
    link_minor_resistance: np.ndarray
    link_type: list[str]
    link_pump_index: np.ndarray
    link_pump_count: int
    link_valve_index: np.ndarray
    pump_coefficients: np.ndarray
    closed_pipe_index: np.ndarray
    node_index: np.ndarray
    junction_index: np.ndarray
    reservoir_index: np.ndarray
    tank_index: np.ndarray
    node_elevations: np.ndarray
    node_name_id: list
    node_x: np.ndarray
    node_y: np.ndarray
    from_node: np.ndarray
    to_node: np.ndarray
    tank_diameter: np.ndarray
    tank_min_level: np.ndarray
    tank_max_level: np.ndarray
    tank_init_level: np.ndarray
    tank_area: np.ndarray
    inp_path: str
    handle: object

```

method n_nodes

```
def n_nodes(self) -> int:
```

method n_links

```
def n_links(self) -> int:
```

`_derive_pump_coefficients`

```
def _derive_pump_coefficients(d, n_pumps: int) -> np.ndarray:
    """Derive pump curve coefficients [h0, r, v=2] from EPANET head curves.

    For a single design point (Qd, Hd) EPANET's convention gives
    'h0 = 1.3333 Hd' and 'r = h0 / (2 Qd)**2' (so H_gain = h0 - r q^2 passes
    through (0, 1.3333 Hd), (Qd, Hd) and (2 Qd, 0)); multi-point curves are fit
    to H = h0 - r Q^v.

    Pumps with no head curve (EPANET curve index 0 -- e.g. constant-power pumps)
    cannot be modeled by this FSP head-gain form; they yield NaN coefficients and
    must be supplied via 'NetworkSpec.pump_coefficients'."""
```

`read_inp`

```
def read_inp(inp_path) -> RawNetwork:
    """Load an EPANET .inp file and extract static OWF data (ports read_inp.m)."""
```

`simulate_with_schedule`

```
def simulate_with_schedule(inp_path, pump_links_lbased, onoff: np.ndarray, time: int, n_nodes: int, n_links: int,
    bypass_links=None) -> tuple[np.ndarray, np.ndarray]:
    """Impose a pump on/off schedule in EPANET and return the true hydraulics.

    Deletes existing controls/rules, drives each pump's status hour-by-hour from
    'onoff' (Pu x T), runs the extended-period simulation, and returns
    (heads, flows) shaped (N x time), (L x time) sampled at hour boundaries.
    Shared by the schedule-imposed validation and the EPANET warm-start.

    'bypass_links' is an optional list of '(link_lbased, pump_pos)' pairs for
    switched bypasses: each is opened exactly when its pump is off. Without this,
    deleting the controls would leave a bypass stuck at its initial status and the
    simulation would be meaningless (e.g. Net3's pipe 330)."""
```

`read_controls_rules`

```
def read_controls_rules(inp_path) -> dict:
    """Return the operational logic EPANET uses, straight from the .inp file:
    the [CONTROLS] (simple) and [RULES] (rule-based) sections as text lines.

    This is the *baseline* schedule the C-OWPF cost is compared against."""
```

`run_epanet`

```
def run_epanet(raw: RawNetwork) -> tuple[np.ndarray, np.ndarray, np.ndarray, np.ndarray]:
    """Extended-period hydraulic simulation (ports init_epanet.m).

    Returns (flows, heads, headloss, demand) each shaped (n_steps, n_links|n_nodes),
    i.e. time along axis 0 -- matching the raw EPANET output ordering the MATLAB
    code transposes downstream.

    Loads a fresh EPANET handle from the source .inp so it stays robust even after
    other simulations (e.g. the warm-start) have opened and unloaded handles."""
```

7.4 owf/connection_matrices.py

Module purpose. Topology / incidence matrices for the WDN (ports ConnectionMatrices_WDN.m).

`class Matrices`

```
class Matrices:
    # fields:
    Pi: np.ndarray
    Pi_telda: np.ndarray
    Pi_prime: np.ndarray
    Pi_reduced: np.ndarray
    Lambda: np.ndarray
    Theta: np.ndarray
```

```

Tau: np.ndarray
Kappa: np.ndarray
Omega: np.ndarray
Delta: np.ndarray
pipe_index: np.ndarray
bypass_index: np.ndarray
Pi_telda_bypass: np.ndarray
Pi_prime_bypass: np.ndarray
Omega_bypass: np.ndarray
S_bypass_pump: np.ndarray

```

build_matrices

```
def build_matrices(raw: RawNetwork, bypass_index: np.ndarray=None, bypass_pump_pos: np.ndarray=None) -> Matrices:
```

7.5 owf/network.py

Module purpose. WDN assembly (ports WDN_setup_IEEE_ACCESS.m + definebounds_WDN.m).

class TankParams

```

class TankParams:
    # fields:
    area: np.ndarray
    init_head: np.ndarray
    min_head: np.ndarray
    max_head: np.ndarray
    mean_head: np.ndarray
    del_tk_tanks: np.ndarray
    Atk: np.ndarray
    Btk: np.ndarray

```

class Bounds

```

class Bounds:
    # fields:
    min_nodal_heads: np.ndarray
    max_nodal_heads: np.ndarray

```

class WDN

```

class WDN:
    # fields:
    config: SolverConfig
    spec: cfg.NetworkSpec
    raw: RawNetwork
    M: Matrices
    pump: PumpParams
    tank: TankParams
    bounds: Bounds
    time: int
    price_final: np.ndarray
    junction_demand_profile: np.ndarray
    lin0: LinPoint
    int_eps: np.ndarray
    int_onoff: np.ndarray
    pump_avail: dict = None

```

method n_nodes

```
def n_nodes(self) -> int:
```

method n_links

```
def n_links(self) -> int:
```

method n_pumps

```
def n_pumps(self) -> int:
```

method n_pipes

```
def n_pipes(self) -> int:
```

method n_tanks

```
def n_tanks(self) -> int:
```

method n_reservoirs

```
def n_reservoirs(self) -> int:
```

method n_junctions

```
def n_junctions(self) -> int:
```

_pump_params

```
def _pump_params(raw: RawNetwork, spec: cfg.NetworkSpec) -> PumpParams:
```

_tank_params

```
def _tank_params(raw: RawNetwork, time: int) -> TankParams:
```

_bounds

```
def _bounds(raw: RawNetwork, tank: TankParams, time: int) -> Bounds:
```

_price

```
def _price(config: SolverConfig, time: int) -> np.ndarray:
```

_resolve_bypasses

```
def _resolve_bypasses(raw: RawNetwork, spec: cfg.NetworkSpec):  
    """Map spec.bypasses {bypass_id: pump_id} to link indices + pump positions."""
```

_resolve_availability

```
def _resolve_availability(raw: RawNetwork, spec: cfg.NetworkSpec) -> dict:  
    """Map spec.pump_availability {pump_id: (start,end)} to {pump_pos: (start,end)}."""
```

setup

```
def setup(config: SolverConfig) -> WDN:  
    """Build the full WDN problem data for the given configuration."""
```

7.6 owf/initial_values.py

Module purpose. Initial linearization point (ports Initial_Values_WDN.m).

initial_point

```
def initial_point(raw: RawNetwork, M: Matrices, pump: PumpParams, bounds, time: int, choice: int, flows_ep: np.  
    ndarray, heads_ep: np.ndarray) -> tuple[LinPoint, np.ndarray, np.ndarray]:  
    """Return (initial LinPoint, stacked iterate int_eps, initial OnOff).  
  
    'flows_ep'/'heads_ep' are EPANET's computed time series (steps x L|N),  
    supplied by the caller so EPANET is only run once."""
```

7.7 owf/linearization.py

Module purpose. Successive-linearization coefficients (ports CalculateNewIterationValues.m).

class PumpParams

```
class PumpParams:
    # fields:
    h0: np.ndarray
    r_m: np.ndarray
    v_m: np.ndarray
    c_m: float
    max_flow: np.ndarray
```

class LinPoint

```
class LinPoint:
    """Frozen linearization coefficients for one MILP solve."""
    # fields:
    Cp: np.ndarray
    C1M: np.ndarray
    C2M: np.ndarray
    APrimePump: np.ndarray
    BPrimePump: np.ndarray
    Cp_bypass: np.ndarray = None
```

linearize

```
def linearize(flows: np.ndarray, M: Matrices, pump: PumpParams) -> LinPoint:
    """Build linearization coefficients around a flow field 'flows' (L x T)."""
```

stack_eps

```
def stack_eps(heads: np.ndarray, flows: np.ndarray, onoff: np.ndarray) -> np.ndarray:
    """Concatenate [Heads; Flows; OnOff] (the successive-linearization iterate)."""
```

error_norm

```
def error_norm(eps: np.ndarray, prev_eps: np.ndarray) -> float:
    """Euclidean norm of the change in the stacked iterate."""
```

7.8 owf/constraints.py

Module purpose. OWF constraints in CVXPY (ports the define*_CVX.m modules, one per function).

reservoir_head

```
def reservoir_head(model, wdn):
    """Theta H == Theta H_min (fix reservoir heads to their elevation)."""
```

tank_flow_aux

```
def tank_flow_aux(model, wdn):
    """TankFlow_aux == Tau' (-Pi) Flows (net inflow to each tank)."""
```

tank_state_space

```
def tank_state_space(model, wdn):
    """Hdummy_i = A_tk H0_i + (del/A_i) Btk TankFlow_aux_i (integrator)."""
```

tank_hdummy_head

```
def tank_hdummy_head(model, wdn):
    """Tau' H == [H0 , Hdummy(:,1:T-1)] (tank head = previous dummy level)."""
```

tank_head_bounds

```
def tank_head_bounds(model, wdn):
    """H_min <= tank heads <= H_max (expressed over Tau)."""
```

pipe_head_loss

```
def pipe_head_loss(model, wdn):  
    """Pi_telda H == Cp + Pi_prime Q (linearized Hazen-Williams)."""
```

mass_balance

```
def mass_balance(model, wdn):  
    """Pi_reduced Q == -demand (conservation of mass at junctions)."""
```

closed_pipes_zero

```
def closed_pipes_zero(model, wdn):  
    """Permanently-closed pipes carry no flow."""
```

switched_bypasses

```
def switched_bypasses(model, wdn):  
    """Bypass pipes that are OPEN iff their pump is OFF.  
  
    Let z = OnOff of the controlling pump (via S). When z=1 (pump on) the bypass  
    is closed: zero flow, head loss relaxed. When z=0 (pump off) the bypass is a  
    normal pipe: flow free, Hazen-Williams head loss enforced."""
```

pump_availability

```
def pump_availability(model, wdn):  
    """Force pumps off outside their source-availability window."""
```

pump_flow

```
def pump_flow(model, wdn):  
    """q_min OnOff <= Lambda Q <= q_max OnOff (flow only when pump is on)."""
```

pump_negative_headloss

```
def pump_negative_headloss(model, wdn):  
    """C1M + C2M (Lambda Q) <= 0."""
```

pump_power

```
def pump_power(model, wdn):  
    """Ppump == A'(Lambda Q) + B' OnOff (linearized FSP power)."""
```

pump_bigm

```
def pump_bigm(model, wdn):  
    """Big-M pump head-gain curve, enforced only when OnOff == 1."""
```

junction_head_bounds

```
def junction_head_bounds(model, wdn):  
    """H_min <= junction heads <= H_max."""
```

tank_terminal

```
def tank_terminal(model, wdn):  
    """Hdummy(:,end) >= tank mean head (terminal level target)."""
```

junction_head_bounds_soft

```
def junction_head_bounds_soft(model, wdn):
```

tank_head_bounds_soft

```
def tank_head_bounds_soft(model, wdn):
```

tank_terminal_soft

```
def tank_terminal_soft(model, wdn):
```

build_constraints

```
def build_constraints(model, wdn, include_mass_balance: bool=True, soft: bool=False):  
    """Assemble the full FSP OWF constraint set (optionally with soft head bounds)."""
```

_energy_cost

```
def _energy_cost(model, wdn):
```

objective

```
def objective(model, wdn):  
    """min sum_t price(t) * sum_pumps(Ppump)/1000    (energy cost)."""
```

objective_soft

```
def objective_soft(model, wdn):  
    """Energy cost plus a penalty on head-bound slacks (penalty CCP warm-start)."""
```

7.9 owf/solver.py

Module purpose. Successive linear-approximation loop (ports WDN_OWf_IEEEACCESS_cvx.m).

class Model

```
class Model:  
    """CVXPY decision variables and linearization parameters."""  
    # fields:  
    Flows: cp.Variable  
    Heads: cp.Variable  
    Ppump: cp.Variable  
    Hdummy: cp.Variable  
    OnOff: cp.Variable  
    TankFlow_aux: cp.Variable  
    Cp: cp.Parameter  
    C1M: cp.Parameter  
    C2M: cp.Parameter  
    APrime: cp.Parameter  
    BPrime: cp.Parameter  
    s_jlo: Optional[cp.Variable] = None  
    s_jhi: Optional[cp.Variable] = None  
    s_tlo: Optional[cp.Variable] = None  
    s_thi: Optional[cp.Variable] = None  
    s_term: Optional[cp.Variable] = None  
    penalty: Optional[cp.Parameter] = None  
    Cp_bypass: Optional[cp.Parameter] = None
```

class OWFResult

```
class OWFResult:  
    # fields:  
    status: str  
    n_iter: int  
    converged: bool  
    objective: float  
    heads: np.ndarray  
    flows: np.ndarray  
    onoff: np.ndarray  
    ppump_linear: np.ndarray  
    ppump_true: np.ndarray  
    errors: list = field(default_factory=list)  
    objectives: list = field(default_factory=list)  
    max_slack: float = float('nan')
```


`_build_model`

```
def _build_model(wdn: WDN) -> Model:
```

`_set_params`

```
def _set_params(model: Model, lin: LinPoint) -> None:
```

`_max_slack`

```
def _max_slack(model: Model) -> float:
```

`_true_pump_power`

```
def _true_pump_power(wdn: WDN, flows: np.ndarray) -> np.ndarray:
    """Nonlinear FSP power  $c_m (h_0 - r f^v) f$  at a flow solution (general  $v$ )."""
```

`solve_owf`

```
def solve_owf(wdn: WDN, lin_override: Optional[LinPoint]=None, eps_override: Optional[np.ndarray]=None) ->
    OWFResult:
    """Run the successive-linearization loop.

    ‘lin_override’ / ‘eps_override’ let a caller seed the initial linearization
    and iterate from an arbitrary warm-start point (e.g. EPANET flows under a
    chosen pump schedule) instead of ‘wdn.lin0’ / ‘wdn.int_eps’."""
```

7.10 `owf/warmstart.py`

Module purpose. EPANET-based multi-start warm-starting for the successive approximation.

`epanet_default_onoff`

```
def epanet_default_onoff(wdn: WDN, thresh: float=0.001) -> np.ndarray:
    """Pump on/off implied by EPANET’s own (rule-based) operation."""
```

`candidate_schedules`

```
def candidate_schedules(wdn: WDN) -> dict[str, np.ndarray]:
    """A small library of pump on/off schedules to warm-start from."""
```

`warmstart_point`

```
def warmstart_point(wdn: WDN, onoff: np.ndarray):
    """Linearization point + stacked iterate from EPANET flows under ‘onoff’."""
```

`_score`

```
def _score(r: OWFResult) -> tuple:
    """Rank results: converged first, then smaller bound violation, then cheaper."""
```

`solve_multistart`

```
def solve_multistart(wdn: WDN, schedules: Optional[dict[str, np.ndarray]]=None, verbose: bool=True) -> tuple[
    OWFResult, str, dict[str, OWFResult]]:
    """Run the successive approximation from each candidate EPANET warm-start.

    Returns (best_result, best_name, all_results). Best = converged, then least
    head-bound violation, then lowest energy cost."""
```

`true_energy_cost`

```
def true_energy_cost(wdn: WDN, flows: np.ndarray) -> float:
    """Energy cost using the TRUE nonlinear FSP power (not the linearized model).

    This is the honest basis for comparing two schedules -- the per-iteration
    objective uses linearized power, which differs slightly between linearization
    points."""
```

`_converge_fixed`

```
def _converge_fixed(wdn: WDN, onoff: np.ndarray, seed_flows: np.ndarray, seed_heads: np.ndarray, max_iter: int
=20) -> OWFResult:
    """Fix a schedule and converge the continuous problem from a given seed."""
```

`_apply_availability`

```
def _apply_availability(wdn: WDN, sched: np.ndarray) -> np.ndarray:
    """Zero out hours where a pump's source is unavailable."""
```

`duty_preserving_schedule`

```
def duty_preserving_schedule(wdn: WDN) -> np.ndarray:
    """Load-shift EPANET's operation: keep each pump's total on-hours but move
    them to the cheapest hours.
```

The price-quantile candidates run pumps **fewer** hours (only the cheapest), which under-pumps and drains the tanks below their minimum. Preserving each pump's EPANET duty and merely re-timing it to cheap hours keeps the tanks in bounds while still cutting the time-of-use bill -- the feasible way to save on high-duty-cycle systems (e.g. BWSN's 11 h / 21 h pumps)."""

`price_threshold_schedules`

```
def price_threshold_schedules(wdn: WDN) -> dict[str, np.ndarray]:
    """Candidate schedules that run pumps only during the cheapest hours.
```

A family of price quantiles (plus all-on and a duty-preserving load-shift) -- physically motivated for time-of-use pricing and cheap to enumerate."""

`optimize_schedule`

```
def optimize_schedule(wdn: WDN, inner_iter: Optional[int]=None, feas_tol: float=2.0, polish: bool=True, max_flips
: int=30, verbose: bool=True) -> tuple[OWFResult, dict]:
    """Optimize the pump on/off schedule and report savings vs EPANET.
```

Each candidate schedule is evaluated **honestly**: fix it, converge the continuous hydraulics, and score the TRUE (nonlinear) energy cost plus the head-bound violation. We do not let the frozen linearization **predict** a distant schedule's quality -- that is what makes a single free-binary MILP **return** the incumbent (its coefficients simply don't apply far away).

Stages:

1. **Baseline** -- EPANET's own operation.
2. **Candidates** -- price-quantile schedules (run only during the cheapest hours), all-on, and a MILP proposal at the incumbent's linearization.
3. **Polish** (optional) -- 1-opt: flip single pump-hours of the winner and keep any flip that lowers cost, giving a schedule that is locally optimal w.r.t. single-hour changes.

Ranking: feasible (slack <= feas_tol) first, then lowest true cost.

Returns (best result, info dict with baseline/best cost, savings and trace)."""

`solve_from_epanet`

```
def solve_from_epanet(wdn: WDN) -> OWFResult:
    """Reproduce EPANET's own operation: seed the linearization from EPANET's
    computed hydraulics, fix that pump schedule, and converge the continuous
    problem. Robust for multi-pump / switched-bypass networks (e.g. Net3) where
    EPANET's controls already give a valid, feasible operating point."""
```

`solve_fixed_schedule`

```
def solve_fixed_schedule(wdn: WDN, onoff: np.ndarray) -> OWFResult:
    """Pin the pump on/off schedule and converge the continuous problem.
```

With the binaries fixed the per-iteration problem is a pure LP, so the successive linearization converges reliably (as on the tree networks)."""

solve_warmstart

```
def solve_warmstart(wdn: WDN, schedules: Optional[dict[str, np.ndarray]]=None, verbose: bool=True) -> tuple[OWFResult, str]:
    """Two-phase warm-start solve for hard networks.

    Phase 1: multi-start over candidate EPANET schedules to pick a good binary
    pump schedule. Phase 2: fix that schedule and converge the continuous problem.
    Returns (result, best_schedule_name)."""
```

7.11 owf/validation.py

Module purpose. Validate the OWF solution against EPANET hydraulics (ports the tail of Main_OW_F_IIEEE_ACCESS.m): error norms on heads, pipe flows and pump flows, plus a true-power comparison.

class ValidationReport

```
class ValidationReport:
    # fields:
    error_heads: float
    error_pipe_flows: float
    error_pump_flows: float
    pump_power_epanet: np.ndarray
    pump_power_opt_true: np.ndarray
    heads_epanet: np.ndarray
    flows_epanet: np.ndarray
```

method summary

```
def summary(self) -> str:
```

validate

```
def validate(wdn: WDN, result: OWFResult) -> ValidationReport:
```

class ScheduleValidationReport

```
class ScheduleValidationReport:
    """Result of imposing the optimized pump schedule back into EPANET
    (reproduces the paper's Fig. 4 hydraulic-feasibility check)."""
    # fields:
    max_abs_head: float
    max_abs_pipe_flow: float
    max_abs_pump_flow: float
    norm_head: float
    norm_pipe_flow: float
    norm_pump_flow: float
    heads_epanet: np.ndarray
    flows_epanet: np.ndarray
```

method summary

```
def summary(self) -> str:
```

validate_schedule

```
def validate_schedule(wdn: WDN, result: OWFResult) -> ScheduleValidationReport:
    """Fix the optimized pump on/off schedule into EPANET, re-simulate the true
    nonlinear hydraulics, and compare against the optimizer's heads/flows.

    Unlike :func:'validate' (which compares against EPANET's own rule-based
    operation), this drives EPANET with the *optimized* schedule, so a good
    linearization should reproduce EPANET closely."""
```

7.12 owf/plots.py

Module purpose. Diagnostic plots: convergence, EPANET-vs-OWF flows/heads, and pump schedules.

`_plt`

```
def _plt():
```

`plot_convergence`

```
def plot_convergence(result: OWFResult):  
    """Successive-linearization error and objective per iteration."""
```

`plot_pump_schedule`

```
def plot_pump_schedule(wdn: WDN, result: OWFResult):  
    """Optimized pump on/off schedule against the electricity price."""
```

`plot_flows`

```
def plot_flows(wdn: WDN, result: OWFResult, report: ScheduleValidationReport, max_pipes: int=6):  
    """Pump flows and a sample of pipe flows: OWF vs EPANET (imposed schedule)."""
```

`plot_heads`

```
def plot_heads(wdn: WDN, result: OWFResult, report: ScheduleValidationReport, max_junctions: int=5):  
    """Tank and junction heads: OWF vs EPANET (imposed schedule)."""
```

`plot_error_summary`

```
def plot_error_summary(wdn: WDN, result: OWFResult, report: ScheduleValidationReport):  
    """Per-time-step max |error| in heads and flows vs EPANET."""
```

`plot_all`

```
def plot_all(wdn: WDN, result: OWFResult, report: Optional[ScheduleValidationReport]=None, outdir: str | Path=  
    'outputs', prefix: Optional[str]=None) -> list[Path]:  
    """Write the full plot set to 'outdir'; returns the written paths."""
```

7.13 owf/netmap.py

Module purpose. Interactive network-map view (Plotly) for the Streamlit UI.

`extract_map_data`

```
def extract_map_data(wdn, result=None) -> dict:  
    """Collect everything the map needs from a WDN (+ optional OWFResult)."""
```

`build_map_figure`

```
def build_map_figure(data: dict, hour: int | None=None):  
    """Plotly figure of the network; colored by solution state at 'hour'."""
```

`build_animated_map_figure`

```
def build_animated_map_figure(data: dict):  
    """Animated map: a Play button steps through the hours, drawing each link's  
    flow as a direction arrow (angle = flow direction, size = |flow|) over the  
    junction-pressure field. Falls back to the static map if no solution."""
```

8 pdn – Power Distribution Network models

8.1 pdn/__init__.py

Module purpose. Power distribution network (PDN) side of C-OWPF.

8.2 pdn/feeders.py

Module purpose. Distribution-feeder data for the power side of C-OWPF.

8.3 pdn/feeders_paper.py

Module purpose. Two SCE feeders transcribed from the paper's Table I (Fig. 5, 47-bus) and Table II (Fig. 6, 56/57-bus).

_build

```
def _build(label, desc, edges, loads, pv_mw, caps_mvar, vbase, sbase, v0_kv, virtual_slack_load=None) -> dict:
```

8.4 pdn/network.py

Module purpose. PDN object: a feeder's linear voltage model plus its DER / limit settings.

class PDN

```
class PDN:
    """A distribution feeder ready for the coupled optimization (all pu)."""
    # fields:
    key: str
    label: str
    model: LinDistModel
    pv_rating: np.ndarray
    solar_profile: np.ndarray
    vmin: float = DEFAULT_VMIN
    vmax: float = DEFAULT_VMAX
    orig_id: list = field(default_factory=list)
```

method N

```
def N(self) -> int:
```

method pv_buses

```
def pv_buses(self) -> np.ndarray:
```

method limit_pv

```
def limit_pv(self, k: int) -> 'PDN':
    """Keep only the 'k' largest-rated PV sites active (zero the rest)."""
```

method solar

```
def solar(self, T: int) -> np.ndarray:
    """Solar availability tiled/truncated to a T-hour horizon."""
```

method build

```
def build(key: str, pv_sizing: float=1.2, vmin: float=DEFAULT_VMIN, vmax: float=DEFAULT_VMAX) -> 'PDN':
    """Construct a PDN from the feeder registry.

    Each PV inverter's apparent-power rating is 'pv_sizing' x the bus's
    nominal active load (the paper places PV on load-hosting buses). Sizing
    strictly to load -- not a fixed per-unit floor -- keeps PV proportional to
    the feeder regardless of its VA base, avoiding artificial reverse-flow
    overvoltage on high-base feeders (e.g. IEEE-33 at 100 MVA)."""
```

8.5 pdn/lindistflow.py

Module purpose. Linear (LinDistFlow / Kekatos) distribution-network voltage model.

class LinDistModel

```
class LinDistModel:
    """Prebuilt affine voltage model for one feeder (all pu)."""
    # fields:
    N: int
    A_tilde: np.ndarray
    F: np.ndarray
    R: np.ndarray
    X: np.ndarray
    V_k: np.ndarray
    r: np.ndarray
    x: np.ndarray
    p_load: np.ndarray
    q_load: np.ndarray
    b_shunt: np.ndarray
    pv_mask: np.ndarray
    v0_sq: float
    parent: np.ndarray
```

method v2

```
def v2(self, p_net: np.ndarray, q_net: np.ndarray) -> np.ndarray:
    """Squared bus voltages for given net loads (pu).

    ‘q_net’ is the net reactive load EXCLUDING shunt caps -- the caps are
    voltage-dependent (q*v) and already folded into ‘R’/‘X’/‘V_k’ via the
    Kekatos (I X*diag(q)) correction, matching paper eq. (1b)/(2b)."""
```

method voltage

```
def voltage(self, p_net: np.ndarray, q_net: np.ndarray) -> np.ndarray:
```

build_lindist

```
def build_lindist(feeder: dict) -> LinDistModel:
    """Assemble the affine voltage model from a ‘pdn.feeders’ entry."""
```

branch_flow_matrix

```
def branch_flow_matrix(model: LinDistModel) -> np.ndarray:
    """Subtree-membership matrix T (N x N): ‘P_branch = T @ p_net’.

Row i (the branch feeding bus i) sums the net injections of bus i and all its
    descendants -- the lossless LinDistFlow branch flow. With per-unit branch flows
    the network loss is ‘sum_i r_i (P_branch_i^2 + Q_branch_i^2)’.


```

8.6 pdn/powerflow.py

Module purpose. Nonlinear single-phase Z-bus (fixed-point) power flow – the **validator**.

zbus_solve

```
def zbus_solve(model: LinDistModel, p_net: np.ndarray, q_net: np.ndarray, tol: float=1e-10, max_iter: int=1000):
    """Full nonlinear Z-bus solve.

    Returns ‘(vmag, vcomplex, loss)’ : non-slack voltage magnitudes (pu), their
    complex values, and the total real network loss (pu) from the exact branch
    currents ‘I_br = y (A_tilde V_full)’ as ‘sum |I_br|^2 r’.

    ‘p_net’ / ‘q_net’ are net loads in pu (positive = consumption) with PV and
    pump loads already folded in, but **excluding** shunt caps -- those are added
    here as a Y-bus susceptance (the physically-correct, voltage-dependent model)."""
```

zbus_powerflow

```
def zbus_powerflow(model: LinDistModel, p_net: np.ndarray, q_net: np.ndarray, tol: float=1e-10, max_iter: int=1000) -> np.ndarray:
    """Nonlinear bus voltage magnitudes (pu) -- thin wrapper over 'zbus_solve'."""
```

8.7 pdn/opf.py

Module purpose. Standalone distribution-network reactive-power OPF.

class PDNOPFResult

```
class PDNOPFResult:
    # fields:
    feeder: str
    q_pv: np.ndarray
    p_pv: np.ndarray
    pv_buses: np.ndarray
    v_lin: np.ndarray
    v_nl: np.ndarray
    v_nl_base: np.ndarray
    loss_nl: np.ndarray
    loss_base: np.ndarray
    p_net: np.ndarray
    q_net: np.ndarray
    v_violation: float
    SBase: float
```

method loss_kw

```
def loss_kw(self) -> np.ndarray:
```

method loss_base_kw

```
def loss_base_kw(self) -> np.ndarray:
```

method loss_reduction_pct

```
def loss_reduction_pct(self) -> float:
```

pump_load_to_bus

```
def pump_load_to_bus(pdn: PDN, ppump_kw: np.ndarray, pump_bus, scale: float=1.0):
    """Map pump electrical power (Pu x T, kW) to a per-bus active load (N x T, pu)."""
```

solve_pdn_opf

```
def solve_pdn_opf(pdn: PDN, T: int, pump_load_pu: np.ndarray | None=None, load_shape: np.ndarray | None=None,
    vmin: float=0.95, vmax: float=1.05, soft_voltage: bool=True, voltage_penalty: float=10000.0, solver: str='
    HIGHS') -> PDNOPFResult:
    """Dispatch PV reactive to hold voltage, then verify with the nonlinear Z-bus."""
```

9 coupled – Coupled C-OWPF solver

9.1 coupled/__init__.py

Module purpose. Coupled Optimal Water-Power Flow (C-OWPF).

9.2 coupled/config.py

Module purpose. Configuration for the coupled Optimal Water-Power Flow (C-OWPF).

class CoupledConfig

```
class CoupledConfig:
    """How a water network and a distribution feeder are co-optimized."""
    # fields:
    feeder: str = 'ieee13'
    pump_bus: Optional[Sequence[int]] = None
    pump_load_scale: float = 1.0
    load_scale: float = 1.0
    load_profile: Optional[np.ndarray] = None
    enable_pv: bool = True
    pv_sizing: float = 1.2
    pump_pf: float = 0.9
    include_loss: bool = True
    vmin: float = 0.95
    vmax: float = 1.05
    soft_voltage: bool = True
    voltage_penalty: float = 10000.0
```

method load_shape

```
def load_shape(self, T: int) -> np.ndarray:
```

9.3 coupled/coupled_lp.py

Module purpose. Coupled Optimal Water-Power Flow (C-OWPF) solver.

class CoupledResult

```
class CoupledResult:
    # fields:
    status: str
    n_iter: int
    converged: bool
    objective: float
    energy_cost: float
    heads: np.ndarray
    flows: np.ndarray
    onoff: np.ndarray
    ppump_true: np.ndarray
    water_max_slack: float
    voltage: np.ndarray
    v2: np.ndarray
    pv_p: np.ndarray
    pv_q: np.ndarray
    pv_buses: np.ndarray
    p_net: np.ndarray
    q_net: np.ndarray
    grid_kw: np.ndarray
    pump_bus: np.ndarray
    v_min: float
    v_violation: float
    loss_kw: np.ndarray = None
    loss_cost: float = float('nan')
    total_cost: float = float('nan')
    errors: list = field(default_factory=list)
    objectives: list = field(default_factory=list)
```

default_pump_buses

```
def default_pump_buses(pdn: PDN, n_pumps: int) -> np.ndarray:
    """Place pumps at the feeder's electrically weakest (lowest nominal-voltage)
    buses, where a pump load most stresses the network -- the interesting case."""
```


`_pv_capability`

```
def _pv_capability(pdn: PDN, cc: CoupledConfig, T: int):
    """PV active output (fixed at availability) and the exact reactive capability.

    PV *active*  $p_{pv}[k,t] = rating[k] * solar(t)$  is fixed (the paper controls PV
    reactive, not active). With active known, the smart-inverter capability circle
     $|q| \leq \sqrt{S^2 - p^2}$  is a CONSTANT bound per (bus, hour) -- exact and LP-safe."""
```

`build_coupled_problem`

```
def build_coupled_problem(wdn: WDN, pdn: PDN, cc: CoupledConfig, trust_region: Optional[tuple]=None):
    """Construct the augmented CVXPY problem; returns (problem, model, ctx).

    ‘‘trust_region=(z0, K)’’ -- only meaningful when the schedule is free (MILP) --
    limits the Hamming distance of the pump on/off variable to at most ‘‘K’’ flips
    from an incumbent schedule ‘‘z0’’, keeping the warm-started linearization valid."""
```

`_feeder_sbase`

```
def _feeder_sbase(pdn: PDN) -> float:
```

`solve_coupled`

```
def solve_coupled(wdn: WDN, pdn: PDN, cc: CoupledConfig, lin_override: Optional[LinPoint]=None, eps_override:
    Optional[np.ndarray]=None, trust_region: Optional[tuple]=None) -> CoupledResult:
    """Run the coupled successive-linearization loop and return a CoupledResult."""
```

`_failed_result`

```
def _failed_result(status, n_iter, ctx, errors, objectives) -> CoupledResult:
    """Return a clearly-failed CoupledResult (e.g. infeasible hard voltage)."""
```

`_snapshot`

```
def _snapshot(wmodel, wdn, pdn, cc, ctx) -> dict:
    """Read current variable values into plain numpy (safe to store)."""
```

9.4 coupled/runner.py

Module purpose. High-level drivers for the coupled water-power optimization.

`setup`

```
def setup(net_num: int, cc: CoupledConfig, time: Optional[int]=None, price_choice: int=1, solver: str='HIGHS') ->
    tuple[WDN, PDN]:
    """Build the water network and the distribution feeder for a coupled run."""
```

`solve_coupled_schedule`

```
def solve_coupled_schedule(wdn: WDN, pdn: PDN, cc: CoupledConfig, onoff: np.ndarray, soft_bounds: bool=True,
    max_iter: int=20) -> CoupledResult:
    """Fix ‘‘onoff’’, warm-start from EPANET, and converge the coupled problem."""
```

`solve_coupled_epanet`

```
def solve_coupled_epanet(wdn: WDN, pdn: PDN, cc: CoupledConfig) -> CoupledResult:
    """Coupled solve reproducing EPANET’s own (rule-based) pump schedule."""
```

9.5 coupled/schedule.py

Module purpose. Coupled pump-schedule optimization (voltage-aware).

`_score`

```
def _score(wdn, res: CoupledResult, feas_tol: float, v_tol: float):
    """Rank key: (infeasible?, total violation, total coupled cost pump+loss)."""
```

`optimize_coupled_schedule`

```
def optimize_coupled_schedule(wdn, pdn, cc, v_tol: float=0.02, feas_tol: float=2.0, inner_iter: int=15, trust_k:
    int | None=None, polish: bool=True, max_flips: int=20, verbose: bool=True, use_milp: bool=True, candidates:
    dict | None=None):
    """Search for the pump schedule minimizing coupled cost + voltage feasibility.

    ‘‘candidates‘‘ overrides the default price-quantile candidate set (pass a small
    dict to cut runtime on big networks); ‘‘use_milp=False‘‘ skips the trust-region
    MILP (the most expensive stage on a large feeder + many-pump network).

    Returns (best CoupledResult, info dict)."""
```

9.6 coupled/validation.py

Module purpose. Validate a coupled solution against BOTH ground-truth simulators.

`coupled_loss_kwh`

```
def coupled_loss_kwh(pdn: PDN, res: CoupledResult) -> float:
    """Total true (nonlinear Z-bus) real network loss over the horizon, in kW*h."""
```

`class CoupledValidationReport`

```
class CoupledValidationReport:
    # fields:
    water: ScheduleValidationReport
    v_lin: np.ndarray
    v_nl: np.ndarray
    v_err_max: float
    v_err_mean: float
    vmin_nl: float
    v_violation_nl: float
```

method summary

```
def summary(self) -> str:
```

`validate_coupled`

```
def validate_coupled(wdn: WDN, pdn: PDN, res: CoupledResult, vmin: float=0.95, vmax: float=1.05) ->
    CoupledValidationReport:
    """Cross-check a coupled result against EPANET (water) and Z-bus (power)."""
```

10 ui – Streamlit application

10.1 ui/___init___py

Module purpose. UI building blocks for the C-OWPF Streamlit app.

10.2 ui/theme.py

Module purpose. Shared look-and-feel for the C-OWPF app: appearance modes, CSS, headers, cards.

`_surface_css`

```
def _surface_css(p: dict) -> str:
    """CSS rules that force one palette on the app surfaces (no media query)."""
```

`_register_plotly`

```
def _register_plotly(mode: str) -> None:
    """Give every Plotly chart a theme-neutral font so axes read on any background."""
```

`inject_theme`

```
def inject_theme(st, mode: str='System') -> None:
    """Inject base CSS + the chosen appearance mode; tune Plotly fonts to match."""
```

`inject_css`

```
def inject_css(st) -> None:
```

`hero`

```
def hero(st, title: str, subtitle: str, tagline: str='') -> None:
```

`section_header`

```
def section_header(st, color: str, title: str, subtitle: str='') -> None:
```

10.3 `ui/landing.py`

Module purpose. Landing page: paper intro, the C-OWPF solve flowchart, a successive-linearization animation with a "why it's better" explainer, the networks, and key formulations.

`_method_note_download`

```
def _method_note_download(st) -> None:
    """Offer the one-page 'how each solve mode is derived in code' transparency note."""
```

`_flowchart_svg`

```
def _flowchart_svg() -> str:
    """Detailed flowchart of how the coupled C-OWPF problem is solved."""
```

`successive_approx_fig`

```
def successive_approx_fig():
    """Animate successive linearization converging on the nonlinear head-loss curve."""
```

`_networks_table`

```
def _networks_table():
```

`render_landing`

```
def render_landing(st) -> None:
```

10.4 `ui/water_section.py`

Module purpose. Water tab: the decoupled FSP Optimal Water Flow (direct / warmstart / optimize / EPANET modes), schedules, EPANET validation, plus per-session comparison and diff.

`_csv_bytes`

```
def _csv_bytes(df: pd.DataFrame) -> bytes:
```

`_build_exports`

```
def _build_exports(case, wdn, result) -> dict:
```

`_session_df`

```
def _session_df(records) -> pd.DataFrame:
```

`_render_case`

```
def _render_case(st, rec) -> None:
```

`_controls`

```
def _controls(st):
```

`render_water`

```
def render_water(st) -> None:
```

10.5 ui/power_section.py

Module purpose. Power tab: standalone distribution-network reactive OPF with DER controls.

`_feeder_choice_labels`

```
def _feeder_choice_labels():
```

`_epanet_pump_kw`

```
def _epanet_pump_kw(net: int):
    """EPANET (rule-based) pump electrical power for a water net.

    Returns '(ppump, ids, Tw)' at the water net's OWN horizon 'Tw' -- the PDN
    OPF then runs over exactly that horizon so pump load and feeder load line up."""
```

`render_power`

```
def render_power(st) -> None:
```

`_render_power_result`

```
def _render_power_result(st) -> None:
```

10.6 ui/coupled_section.py

Module purpose. Coupled tab: joint water-power C-OWPF vs the decoupled hand-off, cost compared.

`_probe_pump_ids`

```
def _probe_pump_ids(net: int):
    """Pump ids for a water net (cached so switching widgets doesn't re-run EPANET)."""
```

`render_coupled`

```
def render_coupled(st) -> None:
```

`_render_coupled_result`

```
def _render_coupled_result(st) -> None:
```

10.7 ui/pdn_plots.py

Module purpose. Plotly figures for the power-distribution side (voltages, PV VAr, loss, map).

feeder_layout

```
def feeder_layout(feeder: dict):
    """Layered left-to-right layout of a radial feeder from its parent array.

    Returns (x, y) over GLOBAL nodes 0..N (0 = slack) and the edge list."""
```

feeder_map

```
def feeder_map(feeder: dict, v_nl: np.ndarray, hour: int, pv_buses, pump_buses, vmin=0.95, vmax=1.05):
    """One-line feeder diagram, buses colored by nonlinear voltage at 'hour'."""
```

voltage_profile

```
def voltage_profile(v_lin: np.ndarray, v_nl: np.ndarray, hour: int, vmin=0.95, vmax=1.05, orig_id=None):
    """Per-bus voltage at one hour: linear (LinDistFlow) vs nonlinear (Z-bus)."""
```

voltage_heatmap

```
def voltage_heatmap(v_nl: np.ndarray, vmin=0.95, vmax=1.05, title='Nonlinear voltage |V| (pu)'):
    """Bus x hour heatmap of the true (Z-bus) voltage magnitude."""
```

pv_reactive_chart

```
def pv_reactive_chart(q_pv: np.ndarray, p_pv: np.ndarray, pv_buses, orig_id=None):
    """PV reactive setpoints (stacked) and total active, over the day."""
```

loss_chart

```
def loss_chart(loss_base_kw: np.ndarray, loss_opt_kw: np.ndarray):
    """True network loss per hour: no-reactive baseline vs optimized setpoints."""
```

11 tests

11.1 tests/test_coupled.py

Module purpose. Coupled water-power (C-OWPF) tests.

test_lindist_matches_nonlinear_nominal

```
def test_lindist_matches_nonlinear_nominal(key):
    """Linear LinDistFlow voltages match the nonlinear Z-bus at nominal load."""
```

test_coupled_fixed_schedule_converges_and_validates

```
def test_coupled_fixed_schedule_converges_and_validates(feeder):
```

test_pv_reactive_supports_voltage

```
def test_pv_reactive_supports_voltage():
    """PV reactive control must lift the minimum voltage vs the no-PV case."""
```

test_coupled_schedule_optimization_saves

```
def test_coupled_schedule_optimization_saves():
    """The voltage-aware schedule search must beat EPANET on cost, stay feasible,
    and its winner must validate against the nonlinear Z-bus."""
```

test_objective_is_pump_cost_only

```
def test_objective_is_pump_cost_only():
    """The reported energy cost must equal the water-only pump cost (feeder-agnostic)."""
```

11.2 tests/test_eightnode.py

Module purpose. Regression tests for the 8-node FSP OWF problem.

solved

```
def solved():
```

test_network_sizes

```
def test_network_sizes():
```

test_converges

```
def test_converges(solved):
```

test_onoff_binary

```
def test_onoff_binary(solved):
```

test_mass_balance_holds

```
def test_mass_balance_holds(solved):
    """Pi_reduced Q == -demand at junctions (a model equality constraint)."""
```

test_reservoir_head_fixed

```
def test_reservoir_head_fixed(solved):
    """Reservoir head is pinned to its elevation."""
```

test_pump_flow_within_bounds

```
def test_pump_flow_within_bounds(solved):
```

test_pump_off_implies_zero_flow

```
def test_pump_off_implies_zero_flow(solved):
    """When OnOff == 0 the pump flow must be (near) zero."""
```

test_power_matches_epanet_aggregate

```
def test_power_matches_epanet_aggregate(solved):
    """Total true FSP power should be in the right ballpark vs EPANET."""
```

11.3 tests/test_more_networks.py

Module purpose. Tests for the additional FSP networks: 3-node (working) and Net1 (loads).

threenode

```
def threenode():
```

test_threenode_sizes

```
def test_threenode_sizes(threenode):
```

test_threenode_pump_coeff_from_epanet

```
def test_threenode_pump_coeff_from_epanet(threenode):
    """Pump coefficients are derived from the EPANET large-pump curve (~533, ~1.3e-6)."""
```

test_threenode_converges

```
def test_threenode_converges(threenode):
```

test_threenode_matches_epanet_schedule

```
def test_threenode_matches_epanet_schedule(threenode):  
    """Optimized schedule re-run in EPANET matches the optimizer closely (Fig. 4)."""
```

test_net1_builds

```
def test_net1_builds():
```

test_net1_default_init_does_not_converge

```
def test_net1_default_init_does_not_converge():  
    """Documents the limitation the warm-start solves."""
```

test_net1_warmstart_converges

```
def test_net1_warmstart_converges():  
    """EPANET multi-start warm-start converges Net1 and matches EPANET (Fig. 4)."""
```

test_net2_builds

```
def test_net2_builds():
```

test_net2_warmstart_converges

```
def test_net2_warmstart_converges():
```

test_ky3_not_registered

```
def test_ky3_not_registered():  
    """KY3 is archived (pumps carry no EPANET head curve) -- not a live network."""
```

test_net3_structure

```
def test_net3_structure():
```

test_net3_reproduces_epanet

```
def test_net3_reproduces_epanet():  
    """solve_from_epanet reproduces EPANET's operation on Net3 (with the bypass)."""
```

11.4 tests/test_optimize_schedule.py

Module purpose. Tests for pump-schedule optimization (savings vs EPANET's own operation).

optimized_net2

```
def optimized_net2():
```

test_price_threshold_family_shape

```
def test_price_threshold_family_shape():
```

test_true_cost_positive

```
def test_true_cost_positive():
```

test_net2_optimization_saves

```
def test_net2_optimization_saves(optimized_net2):  
    """The optimizer should beat EPANET's operation on Net2, feasibly."""
```

test_net2_optimized_schedule_feasible_in_epanet

```
def test_net2_optimized_schedule_feasible_in_epanet(optimized_net2):  
    """Re-simulating the optimized schedule in EPANET gives positive pressures  
    and in-bounds tanks -- i.e. the savings are physically real."""
```

test_availability_respected

```
def test_availability_respected():  
    """Net3's Lake pump must stay off outside its availability window."""
```

11.5 tests/test_plots.py

Module purpose. Smoke tests for the diagnostic plots.

solved

```
def solved():
```

test_individual_plots

```
def test_individual_plots(solved):
```

test_plot_all_writes_files

```
def test_plot_all_writes_files(solved, tmp_path):
```